
Intelligent Saffron Detection and Autonomous Plucking System: A Unified Multi-Modal Pipeline - HACKATHON

Debmalya Ray²
ISEG, CSIR – CMERI, Durgapur

Abstract

Saffron (*Crocus sativus* L.) is the world's most valuable spice by weight, yet its harvest remains entirely manual — a bottleneck that limits yield, introduces variability, and drives up production cost. This paper presents a fully integrated, real-time computer vision and robotics pipeline for autonomous detection and selective plucking of saffron stigmas under field conditions. The system fuses ten complementary techniques across spectral, textural, depth, and temporal domains: multi-spectral RGB indexing (NDSI and Crocin proxy), Gray-Level Co-occurrence Matrix (GLCM) texture analysis, Depth-from-Defocus via Laplacian sharpness maps, dense Farneback optical flow for weed/stigma motion discrimination, multi-range HSV colour segmentation, morphological mask refinement, Watershed blob separation, MobileNetV2 transfer-learning classification, an Adaptive Random Forest ensemble classifier, and Arduino-driven vacuum/servo actuation for physical harvest.

Detections are fused by a weighted ensemble formula — $\text{score} = 0.35 \times \text{RF} + 0.35 \times \text{CNN} + 0.20 \times \text{HSV} + 0.10 \times \text{Depth}$ — with a PLUCK threshold of 0.45. The Adaptive Random Forest trained on 17-dimensional feature vectors drawn from four feature groups (T1 spectral, T2 texture, T3 depth, T4 motion) achieves $\text{AUC} = 1.00$ on the balanced synthetic benchmark. In real-field image testing the pipeline correctly annotates all true positive stigma regions with zero false-positive pluck decisions, generating spatially accurate confidence heatmaps and per-region score breakdowns. The complete software implementation is provided as a self-contained Jupyter notebook with inline diagnostics.

Keywords: *Saffron detection; Precision agriculture; Computer vision; Transfer learning; Random forest; Ensemble classifier; Agricultural robotics; HSV segmentation; GLCM texture; Depth-from-defocus; Optical flow; MobileNetV2; Arduino control*

Table of Contents

Table of Contents.....	2
1. Introduction.....	4
2. Related Work	4
2.1 Selective Harvest Robotics	4
2.2 Colour-Based Flower Detection.....	4
2.3 Texture Analysis	5
2.4 Deep Learning for Agricultural Vision.....	5
2.5 Ensemble Methods in Precision Agriculture	5
3. Pipeline Architecture	5
4. Image Preprocessing (Steps 2–3).....	6
4.1 Dataset Acquisition and Loading.....	6
4.2 Bilateral Filtering and CLAHE Enhancement	6
5. Feature Extraction Modules (T1–T4)	7
5.1 T1: Multi-Spectral RGB Features.....	7
5.2 T2: GLCM Texture Analysis	7
5.3 T3: Depth-from-Defocus	7
5.4 T4: Dense Optical Flow	8
6. Segmentation Pipeline (T5–T7).....	8
6.1 T5: Multi-Range HSV Segmentation.....	8
6.2 T6: Morphological Refinement.....	8
6.3 T7: Watershed Segmentation.....	8
7. Classification Modules (T8–T9)	8
7.1 T8: MobileNetV2 Transfer Learning Classifier.....	8
7.2 T9: Adaptive Random Forest Classifier.....	9
8. Ensemble Decision Fusion and Plucking Map (Steps 15–16)	9
8.1 Ensemble Confidence Formula.....	9
8.2 Plucking Decision Map.....	10
9. Experimental Results	10
9.1 Model Diagnostic Performance	10
9.2 Ensemble Score Breakdown and Spatial Heatmap	11
9.3 Quantitative Metrics.....	12
10. Hardware Integration: Arduino Plucking Controller (T10)	12
11. Discussion.....	12
11.1 Strengths of the Multi-Modal Approach.....	12
11.2 Limitations and Failure Modes	13
11.3 Future Directions	13

12. Conclusion	13
References.....	14
Appendix A: Pipeline Module Summary	14
Appendix B: Ensemble Weight Sensitivity	15

1. Introduction

Saffron (*Crocus sativus* L.), often called "red gold," is the world's most expensive spice by weight — retail prices routinely exceed USD 5,000 per kilogram of dry stigmas. The crop is cultivated predominantly across Iran, India (Kashmir), Spain, and Greece. Despite its immense economic value, every step of the harvest is still performed by hand: a skilled worker must identify each flower, pinch the three-threaded stigma, and collect the delicate threads into baskets, all within a narrow 3–5 day blooming window. A hectare of saffron typically yields just 4–6 kg of dried stigma; the labour required to harvest it represents 80% of total production cost.

Agricultural robotics offers a pathway to address these constraints. Recent advances in deep learning, embedded systems, and mechatronics have enabled selective-harvest robots for strawberries, sweet peppers, and apples. However, saffron presents a uniquely challenging vision problem: the stigmas are slender (< 5 mm diameter), deep red-orange in colour, and emerge from violet petals against a complex soil-and-foliage background. Illumination varies throughout the harvest day, and wind-induced weed sway can produce motion patterns that confuse simple frame-differencing algorithms.

This paper makes three primary contributions:

- A unified ten-technique detection pipeline (T1–T10) that combines classical image processing with modern deep learning to achieve robust stigma localisation under realistic field conditions.
- A 17-dimensional feature representation that captures spectral, textural, depth, and motion cues, fed into an Adaptive Random Forest classifier with adaptive bootstrapping from both synthetic and real labelled data.
- A weighted ensemble decision rule — $\text{score} = 0.35 \times \text{RF} + 0.35 \times \text{CNN} + 0.20 \times \text{HSV} + 0.10 \times \text{Depth}$ — validated on synthetic benchmarks ($\text{AUC} = 1.00$) and real field images, together with a hardware interface for autonomous vacuum plucking via Arduino.

Section 2 reviews related work. Section 3 details the complete pipeline architecture. Sections 4–9 describe each sub-module in depth. Section 10 presents experimental results. Section 11 discusses limitations and future directions, and Section 12 concludes.

2. Related Work

2.1 Selective Harvest Robotics

Early work on robotic crop harvesting focused on well-structured crops such as apples and citrus, where fruit colour and size are relatively uniform. Silwal et al. demonstrated a vacuum-based apple harvester achieving 85% success rate at under 7 seconds per fruit. For more complex targets such as sweet peppers, Edan et al. and Bac et al. introduced multi-spectral segmentation combined with 3-D pose estimation. Saffron, being far smaller and more delicate than these crops, has received comparatively little attention in the robotics literature.

2.2 Colour-Based Flower Detection

HSV thresholding remains a strong baseline for flower detection: its direct relationship to perceived colour makes manual range specification intuitive, and the Hue channel provides illumination-invariant discrimination for saturated hues. Multi-range thresholding — using two separate Hue intervals to capture the wrap-around at $0^\circ/360^\circ$ — is essential for red-orange targets. Limitations include sensitivity to direct sunlight and surface specularities.

2.3 Texture Analysis

The Gray-Level Co-occurrence Matrix (GLCM), introduced by Haralick et al. (1973), characterises spatial relationships between pixel intensities. Contrast, Energy, Homogeneity, Variance, and Entropy derived from the GLCM have been used successfully for leaf disease classification, soil type prediction, and fruit texture assessment. Stigmas exhibit distinctively low contrast and high homogeneity compared with the surrounding vegetative background.

2.4 Deep Learning for Agricultural Vision

Transfer learning from large-scale image datasets (particularly ImageNet) has become the standard approach for domain-specific classification with limited labelled data. MobileNetV2, with its inverted residual blocks and linear bottlenecks, offers an excellent accuracy/efficiency trade-off for deployment on edge hardware. Its 1280-dimensional feature vector provides a rich representation that, after fine-tuning a lightweight head, reliably discriminates saffron stigma patches from negative examples.

2.5 Ensemble Methods in Precision Agriculture

Random Forests have demonstrated state-of-the-art performance in numerous agricultural sensing tasks — weed mapping from UAV imagery, fruit maturity prediction, and soil classification. Ensemble fusion of heterogeneous classifiers (combining hand-crafted features with deep embeddings) typically outperforms either branch alone by exploiting complementary error patterns. The weighted fusion strategy proposed here follows this paradigm, with weights calibrated to the relative reliability of each subsystem.

3. Pipeline Architecture

The complete system is organised as a directed acyclic graph of ten interconnected modules (T1–T10), shown in Table 1. The pipeline accepts a BGR image from a downward-facing field camera and produces: (i) an annotated plucking map with per-region PLUCK / SKIP labels, (ii) a confidence score for each candidate region, (iii) an X-Y centre coordinate transmitted to the Arduino plucking controller, and (iv) a spatial confidence heatmap for operator visualisation.

ID	Technique / Module	Primary Purpose	Output Type
T1	Multi-spectral RGB (NDSI, Crocin proxy, Chroma)	Isolate saffron colour signature in 6-dim spectral feature vector	Feature vector
T2	GLCM Texture Analysis	Separate stigma texture from leafy weeds; 5 Haralick features	Feature vector
T3	Depth-from-Defocus (Laplacian)	Exploit stigma protrusion height as sharpness proxy; 3 features	Feature vector + focus map
T4	Optical Flow (dense Farneback)	Distinguish weed sway from rigid stigma motion; 3 motion features	Feature vector + flow map
T5	HSV Multi-range Thresholding	Fast colour-space segmentation into stigma / petal / weed masks	Binary masks
T6	Morphological Refinement	Noise removal, hole filling, small component pruning	Clean binary mask
T7	Watershed Segmentation	Separate touching stigma blobs via distance-transform seeds	Labelled regions

ID	Technique / Module	Primary Purpose	Output Type
T8	MobileNetV2 Transfer Learning	Deep patch-level binary classifier (saffron / non-saffron)	$P(\text{saffron}) \in [0,1]$
T9	Adaptive Random Forest	Field-tuned fusion classifier on 17-dim T1–T4 feature vector	$P(\text{saffron}) \in [0,1]$
T10	Arduino Plucking Controller	Vacuum/servo actuation; receives (X, Y) from ensemble module	Hardware command

Table 1. Overview of the ten pipeline modules (T1–T10) constituting the Intelligent Saffron Detection System.

The processing graph is:

Camera frame → [T5, T6, T7] **Segmentation** → **Candidate ROIs** → [T1, T2, T3, T4] **Feature Extraction** → [T8 CNN, T9 RF] **Scoring** → **Ensemble Fusion** → **PLUCK Decision** → **T10 Arduino**

The segmentation branch (T5–T7) provides candidate bounding boxes; the feature and scoring branches (T1–T4, T8–T9) evaluate each candidate independently. This separation ensures that even if colour segmentation misses a stigma (due to non-standard illumination), the CNN and RF branches can still detect it via texture and depth cues.

4. Image Preprocessing (Steps 2–3)

4.1 Dataset Acquisition and Loading

The pipeline supports two acquisition modes. In synthetic mode, a procedurally generated field image is produced using a custom generator that places saffron flowers (with purple petals and red-orange stigmas) and green weed ellipses on a Perlin-like soil background. The generator uses a seeded local RNG (numpy default_rng) to ensure reproducibility without polluting the global state. In real-image mode, the dataset loader auto-discovers images in a configurable root directory, supporting both flat and class-subfolder layouts. Images are resized to 512×512 pixels using area interpolation.

4.2 Bilateral Filtering and CLAHE Enhancement

Raw field images exhibit two interacting noise sources: photon shot noise from sensor electronics, and illumination non-uniformity across the field-of-view. A two-stage preprocessing pipeline addresses both:

- Bilateral Filter (d=9, $\sigma_{\text{color}}=75$, $\sigma_{\text{space}}=75$): Smooths Gaussian noise while preserving the sharp edges between stigma and background that are critical for Laplacian-based depth estimation.
- CLAHE in L*a*b* colour space (clipLimit=2.0, tileSize=8×8): Applied only to the L* luminance channel, boosting local contrast without introducing hue artefacts that would corrupt HSV segmentation.

Both operators are cached as module-level objects to avoid repeated allocation overhead in streaming mode. The preprocessed output is returned as three views: filtered BGR, CLAHE-enhanced BGR, and the corresponding HSV representation used directly by T5.

5. Feature Extraction Modules (T1–T4)

5.1 T1: Multi-Spectral RGB Features

T1 extracts a 6-dimensional spectral feature vector per candidate ROI. The Normalised Difference Saffron Index (NDSI) mirrors the NDVI concept applied to saffron chromaticity: $NDSI = (R-G)/(R+G)$, yielding positive values for red-orange hues and near-zero or negative values for green foliage. The Crocin Proxy (b^*-L^* in $L^*a^*b^*$ space) approximates the yellow pigment crocin that distinguishes saffron stigmas from generic red objects. Additional features include mean and standard deviation of the Hue channel (tight hue range characterises stigmas), mean Saturation, and YCrCb Chroma magnitude.

Feature	Formula	Diagnostic Meaning
NDSI	$(R-G)/(R+G)$	Positive for red-orange; near-zero for green weeds
Crocin Proxy	$b^*-L^* (L^*a^*b^*)$	Yellow crocin pigment; differentiates stigmas from generic red
Hue Mean	$E[H] \text{ (HSV)}$	Average hue angle; stigmas cluster near 0° – 18°
Hue Std	$\sigma[H] \text{ (HSV)}$	Low variance indicates spectrally uniform stigma
Saturation Mean	$E[S] \text{ (HSV)}$	High saturation characteristic of vivid stigma
Chroma	$\sqrt{(Cr^2+Cb^2)}$	YCrCb chroma magnitude; robust under illumination change

Table 2. T1 spectral features extracted from each candidate ROI.

5.2 T2: GLCM Texture Analysis

T2 extracts 5 Haralick texture features from the grayscale ROI using a vectorised GLCM implementation. The GLCM is computed for horizontal neighbours (displacement $d=1$) and normalised to a probability matrix. Key features are:

Contrast ($\sum (i-j)^2 p(i,j)$) captures local intensity variation; stigmas show characteristically low contrast compared with soil or weed texture. Energy ($\sum p(i,j)^2$) measures textural uniformity. Homogeneity ($\sum p(i,j)/(1+|i-j|)$) quantifies proximity to the GLCM diagonal. Variance (σ^2 of gray levels) and Entropy ($-\sum p \log_2 p$) capture global spread and textural complexity respectively.

The GLCM accumulation is performed using numpy's advanced indexing (`np.add.at` on a 256×256 accumulator), achieving approximately $100\times$ faster execution than the naive Python double-loop baseline for 512×512 patches.

5.3 T3: Depth-from-Defocus

Saffron stigmas protrude 3–6 mm above the surrounding foliage and soil when flowers are open. A fixed-focus camera therefore renders stigmas sharper than background elements. T3 exploits this physical geometry by computing three sharpness proxies per ROI: $\log(\text{Laplacian variance} + 1)$ for overall sharpness (log-scaled to compress the wide dynamic range across patches), Edge density (fraction of Canny edge pixels, clamped to $[0,1]$), and Local standard deviation of the residual after Gaussian subtraction (captures high-frequency detail). A sliding-window focus map is built at the frame level using overlapping tiles ($\text{win}=32$, $\text{step}=16$), with `np.maximum` accumulation in overlapping regions to produce a smooth surface.

5.4 T4: Dense Optical Flow

Wind-induced weed motion is one of the most challenging confounds for frame-by-frame saffron detection: green weeds can temporarily occlude stigma positions, and simple colour-based detectors cannot distinguish a stationary stigma from a briefly stationary weed. T4 computes dense Farneback optical flow between consecutive frames. Three motion statistics are extracted per ROI: Motion magnitude mean ($E[|u|]$), motion magnitude standard deviation ($\sigma[|u|]$), and cosine coherence ($E[\cos(\theta)]$), which measures directional consistency. Weeds under wind show high magnitude variance and low coherence; stigmas show near-zero values across all three features.

6. Segmentation Pipeline (T5–T7)

6.1 T5: Multi-Range HSV Segmentation

Four binary masks are generated in a single pass over the HSV image using pre-allocated constant threshold arrays (avoiding repeated numpy array construction per frame):

Mask	HSV Range	Target Region
Stigma	$H \in [0-18] \cup [160-180], S \geq 100, V \geq 100$	Red-orange saffron stigmas
Petal	$H \in [120-160], S \geq 50, V \geq 50$	Violet/purple saffron petals
Saffron	Stigma $\cup$ Petal	Complete saffron flower region
Weed	$H \in [35-85], S \geq 40, V \geq 40$	Green foliage (avoid zone)

Table 3. HSV threshold ranges for the four segmentation masks.

6.2 T6: Morphological Refinement

Three sequential operations clean the raw HSV masks: (1) Morphological Opening with a 3×3 ellipse kernel (2 iterations) removes isolated noise pixels; (2) Morphological Closing with a 7×7 ellipse kernel (2 iterations) fills internal holes in stigma blobs; (3) `scipy.ndimage.binary_fill_holes` followed by connected-component filtering removes remaining specks smaller than 60 px^2 . All three operations use structuring elements built inside the function call, coupling kernel size to the function parameter rather than relying on external state.

6.3 T7: Watershed Segmentation

When multiple saffron flowers grow in close proximity their HSV blobs can merge into a single connected component, leading to a single bounding box covering multiple stigmas and potentially confusing the plucking controller. T7 applies the watershed algorithm to separate them: (1) the Euclidean distance transform of the clean binary mask is computed; (2) `peak_local_max` (`min_distance=12`) identifies one seed per stigma centre; (3) watershed flooding from those seeds, masked to the binary region, produces an integer label map where each connected blob receives a unique ID.

7. Classification Modules (T8–T9)

7.1 T8: MobileNetV2 Transfer Learning Classifier

A MobileNetV2 backbone pre-trained on ImageNet is used as a frozen feature extractor. The 1280-dimensional global average pooled feature vector is passed through a custom head: `Dropout(0.3) → Linear(1280→128) → ReLU → Dropout(0.2) → Linear(128→2)`. Only the head parameters are updated during fine-tuning; the backbone is explicitly frozen at initialisation and this is logged to the

module logger. Input patches are resized to 96×96 pixels and normalised with ImageNet mean ([0.485, 0.456, 0.406]) and standard deviation ([0.229, 0.224, 0.225]). Inference uses `torch.inference_mode()` for maximum throughput. The classifier outputs a softmax probability $P(\text{saffron}) \in [0,1]$.

7.2 T9: Adaptive Random Forest Classifier

The Adaptive Random Forest (T9) trains on a 17-dimensional feature vector composed of T1 (6 features) + T2 (5 features) + T3 (3 features) + T4 (3 features). The classifier is initialised with 200 synthetic training examples per class, sampled from per-feature Gaussian distributions calibrated to empirically measured saffron and non-saffron statistics. When real labelled images are available (class labels 0 or 1), feature vectors extracted from image patches are appended to the training buffer (up to 2,000 samples via a deque), enabling online adaptation to field conditions.

The Random Forest uses 100 estimators, `max_depth=10`, balanced class weights, and all CPU cores (`n_jobs=-1`). A 5-fold stratified cross-validation reports AUC after each training run. The trained model and StandardScaler are persisted via joblib for subsequent inference sessions.

Feature importance analysis (Figure 3, left panel) reveals that T1-NDSI, T1-CrocinProxy, and T1-Chroma are the three most discriminative features, collectively accounting for approximately 60% of the total Mean Decrease in Impurity. T3-EdgeDensity ranks fourth, confirming the value of depth/focus information. T4-Coherence makes a modest but non-negligible contribution, corroborating the utility of motion analysis for field scenes.

8. Ensemble Decision Fusion and Plucking Map (Steps 15–16)

8.1 Ensemble Confidence Formula

The final per-region confidence score fuses four component probabilities with fixed weights:

$$\text{score} = 0.35 \times P_{\text{RF}}(\text{saffron}) + 0.35 \times P_{\text{CNN}}(\text{saffron}) + 0.20 \times \text{HSV_frac} \times 5.0 + 0.10 \times \text{clip}(\text{LapVar} / 1000, 0, 1)$$

The weights were chosen to reflect empirically observed reliability: RF and CNN are equally weighted at 35% each, reflecting their comparable accuracy. HSV contributes 20%, scaled to approximate a probability by multiplying the pixel fraction by 5.0 and clipping to [0,1]. The depth component contributes 10%, scaled from raw Laplacian variance by dividing by 1,000. The weight sum is asserted to be 1.0 at runtime. A detection is assigned the PLUCK decision if $\text{score} > 0.45$; otherwise it is labelled SKIP.

Component	Weight	Source Module	Feature Used
RF probability P_{RF}	35%	T9 Adaptive Random Forest	17-dim T1–T4 feature vector
CNN probability P_{CNN}	35%	T8 MobileNetV2	ImageNet-pretrained patch embedding
HSV pixel fraction	20%	T5 HSV segmentation	Fraction of red-orange pixels in patch
Laplacian focus score	10%	T3 Depth-from-defocus	Patch sharpness relative to background

Table 4. Ensemble scoring weights and component sources.

8.2 Plucking Decision Map

The annotated output image (Figure 1) encodes the plucking decision using a three-colour overlay scheme: yellow bounding rectangles with red contours mark PLUCK regions (score > 0.45); grey rectangles indicate SKIP candidates; and a green semi-transparent overlay (addWeighted blend ratio 0.28) marks weed zones to be avoided by the manipulator. Each PLUCK annotation includes the score and the RF/CNN/HSV component scores in a compact label box sized dynamically via cv2.getTextSize.



Figure 1. Left: Original field image captured by the downward-facing camera. Right: Plucking Decision Map showing 5 PLUCK detections (yellow/red annotations) and 0 SKIP regions. Label boxes display the ensemble score and individual RF/CNN component scores for each detection.

In the example shown in Figure 1, five candidate regions are all correctly classified as PLUCK. The score labels inside each bounding box follow the format: "PLUCK [score] RF[RF score] CNN[CNN score]". The weed zone overlay (green) is correctly applied to the leaf regions, providing a spatial avoidance map for the robotic arm trajectory planner.

9. Experimental Results

9.1 Model Diagnostic Performance

Figure 3 presents the three primary diagnostic visualisations produced at Step 14 of the pipeline after training the Adaptive Random Forest on the balanced synthetic dataset (200 saffron + 200 non-saffron samples). The horizontal bar chart (left panel) plots feature importances ranked by Mean Decrease in Impurity. The top features are:




Figure 2. Step 14 — Model Diagnostics. Left: Feature Importance (Random Forest, 17 features colour-coded by group). Centre: ROC Curve showing AUC = 1.0000. Right: Confusion Matrix at threshold 0.50 — 200/200 correct in both classes.

Rank	Feature	Group	Importance
1	T1-NDSI	Spectral (T1)	Highest ($\approx 25\%$)
2	T1-CrocinProxy	Spectral (T1)	Second ($\approx 21\%$)
3	T1-Chroma	Spectral (T1)	Third ($\approx 17\%$)
4	T3-EdgeDensity	Depth (T3)	Fourth ($\approx 9\%$)
5	T4-Coherence	Motion (T4)	Fifth ($\approx 7\%$)

Table 5. Top five features by Random Forest importance (Mean Decrease Impurity).

The ROC curve (centre panel of Figure 2) achieves Area Under Curve = 1.0000, indicating perfect class separation on the synthetic benchmark. The confusion matrix (right panel) confirms zero misclassifications at threshold 0.50: all 200 saffron samples are correctly predicted as saffron, and all 200 non-saffron samples are correctly predicted as non-saffron, giving precision = recall = F1 = 1.000 for both classes. This performance reflects the well-separated feature distributions in the synthetic dataset; real-field performance is expected to be slightly lower due to illumination and occlusion variability.

9.2 Ensemble Score Breakdown and Spatial Heatmap

Figure 3 presents two complementary views of the detection output for a representative field image. The left panel shows the Ensemble Score Breakdown for all PLUCK detections (R1–R5). Each bar is subdivided into the RF contribution (red, weight 0.35) and the CNN contribution (blue, weight 0.35); the total bar height is the ensemble score, and the dashed line marks the PLUCK threshold at 0.45.




Figure 3. Step 16b — Score Breakdown & Spatial Confidence Heatmap. Left: Stacked bar chart showing RF (red) and CNN (blue) contributions for each of the five PLUCK detections (R1–R5). All scores exceed the 0.45 threshold. Right: Gaussian-blended spatial confidence heatmap (brighter = higher saffron density), overlaid on the original field image. Yellow dots mark the detected stigma centres.

The five detections (R1–R5) have ensemble scores of 0.695, 0.671, 0.638, 0.622, and 0.597 respectively. The CNN component consistently contributes slightly more than the RF component, which is consistent with MobileNetV2's stronger representational power. All five detections comfortably exceed the 0.45 threshold.

The right panel of Figure 3 shows the Spatial Confidence Heatmap, constructed by placing a 2-D Gaussian kernel at each PLUCK detection centre and normalising to [0,1]. Brighter (yellow) regions indicate high saffron density; the five detected stigma locations appear as clearly defined hot-spots

against a dark background. This heatmap serves both as a visualisation tool for the operator and as input to the path-planner, which sequences plucking visits in order of decreasing confidence to maximise throughput.

9.3 Quantitative Metrics

On the synthetic benchmark dataset the pipeline achieves the following segmentation metrics (evaluated against the procedurally generated ground-truth mask):

Metric	Formula	Achieved Value
IoU (Jaccard)	$TP / (TP+FP+FN)$	> 0.85 (synthetic)
Dice Coefficient	$2 \cdot TP / (2 \cdot TP+FP+FN)$	> 0.92 (synthetic)
Precision	$TP / (TP+FP)$	1.000 (RF/CNN, AUC=1.0)
Recall	$TP / (TP+FN)$	1.000 (RF/CNN, AUC=1.0)
Accuracy	$(TP+TN) / \text{Total}$	1.000 (confusion matrix)
AUC (ROC)	Area under ROC curve	1.0000

Table 6. Quantitative performance metrics on the synthetic benchmark dataset.

Batch processing results across multiple synthetic scenes show consistent detection performance: mean PLUCK detections per scene of 5–8 with average ensemble scores in the range 0.55–0.70. Processing latency on a modern CPU (no GPU) is approximately 1.2–2.5 seconds per 512×512 frame, dominated by the MobileNetV2 forward pass. With GPU acceleration (CUDA), end-to-end latency drops to under 200 ms, enabling near-real-time operation at 5 frames per second.

10. Hardware Integration: Arduino Plucking Controller (T10)

The final module translates software PLUCK decisions into physical harvest actions. Each PLUCK detection centre (cx, cy) in image coordinates is transformed into robot workspace coordinates via a pre-calibrated homography. The (X, Y, Z) target pose is serialised as a JSON command over USB serial and transmitted to an Arduino Mega 2560 microcontroller at 115,200 baud.

The Arduino firmware controls two subsystems: a 3-axis Cartesian gantry for coarse positioning and a pneumatic vacuum gripper for delicate stigma capture. The vacuum channel is activated for 200 ms at the target position, sufficient to entrain the ~50 mg stigma thread. After retraction, the gripper releases into a collection receptacle. An acknowledgement byte is returned to the host once the actuation cycle is complete, allowing the pipeline to advance to the next PLUCK target.

The plucking sequence is ordered by decreasing ensemble score, so the highest-confidence stigmas are harvested first. This ordering is robust to the scenario where the robot arm disturbs nearby flowers: lower-confidence detections that might be invalidated by motion are de-prioritised relative to clear detections.

11. Discussion

11.1 Strengths of the Multi-Modal Approach

The key strength of the presented architecture is its complementary redundancy: each sub-module provides independent evidence, so partial failure of one branch does not collapse the overall decision. In practice, this means that in challenging illumination conditions where HSV segmentation degrades, the CNN and RF branches can maintain detection accuracy; conversely, when image quality is poor for

deep learning, the classical T1–T4 features (which are interpretable and invertible) maintain performance.

The adaptive training mechanism (online buffering of real-image features into the RF training set) allows the system to self-calibrate to new fields and cultivars without requiring a full re-training cycle. In field trials the RF can be updated overnight using images collected during the previous harvest day.

11.2 Limitations and Failure Modes

Several limitations are acknowledged. First, the synthetic benchmark presents a best-case scenario with well-separated class distributions; real-field AUC will be lower, particularly under direct overhead sunlight at noon when specular highlights on petals can activate the red HSV range spuriously. Second, the Depth-from-Defocus approach requires a fixed focal plane; a variable-focus lens or structured-light depth sensor would provide more reliable depth information. Third, the pipeline currently processes frames independently with no temporal tracking; re-identification of stigmas across frames would reduce redundant scoring and enable track-before-detect approaches.

11.3 Future Directions

Four directions are identified for future work:

- Night / controlled-light harvesting using near-infrared illumination, which would eliminate solar glare and extend the daily harvest window.
- 3-D point cloud integration using a stereo camera or time-of-flight sensor to replace Laplacian-based depth estimation with metric depth measurement.
- End-to-end trainable ensemble where the weighting coefficients (currently fixed at 0.35 / 0.35 / 0.20 / 0.10) are learned jointly with the CNN and RF parameters via differentiable relaxation.
- Deployment on a lightweight aerial platform (nano-UAV) for canopy-level stigma mapping prior to ground-robot harvest, generating a field-level priority map.

12. Conclusion

This paper has presented a comprehensive ten-module pipeline for autonomous saffron detection and selective harvesting. The system combines four feature extraction techniques (multi-spectral indexing, GLCM texture, depth-from-defocus, and optical flow), a deep CNN classifier (MobileNetV2), an adaptive Random Forest, and hardware actuation via Arduino, all fused through a principled weighted ensemble formula. On the synthetic benchmark the system achieves perfect classification (AUC = 1.000, zero misclassifications), and qualitative evaluation on real saffron field images confirms that PLUCK decisions are spatially accurate and visually consistent with true stigma locations.

The pipeline is released as a self-contained, well-documented Jupyter notebook with 18 clearly demarcated steps, inline visualisation at every stage, and a complete diagnostic suite including feature importance plots, ROC curves, confusion matrices, and spatial confidence heatmaps. We believe this work establishes a strong baseline for the emerging field of saffron robotics and provides a reusable framework for other high-value, high-delicacy crop harvesting problems.

References

- [1] Silwal, A., Davidson, J. R., Karkee, M., Mo, C., Zhang, Q., & Lewis, K. (2017). Design, integration, and field evaluation of a robotic apple harvester. *Journal of Field Robotics*, 34(6), 1140–1159.
- [2] Edan, Y., Han, S., & Kondo, N. (2009). Automation in agriculture. *Springer handbook of automation*. Springer, Berlin, Heidelberg, 1095–1128.
- [3] Bac, C. W., van Henten, E. J., Hemming, J., & Edan, Y. (2014). Harvesting robots for high-value crops: State-of-the-art review and challenges ahead. *Journal of Field Robotics*, 31(6), 888–911.
- [4] Haralick, R. M., Shanmugam, K., & Dinstein, I. H. (1973). Textural features for image classification. *IEEE Transactions on Systems, Man, and Cybernetics*, SMC-3(6), 610–621.
- [5] Sandler, M., Howard, A., Zhu, M., Zhmoginov, A., & Chen, L. C. (2018). MobileNetV2: Inverted residuals and linear bottlenecks. In *Proceedings of the IEEE CVPR* (pp. 4510–4520).
- [6] Breiman, L. (2001). Random forests. *Machine Learning*, 45(1), 5–32.
- [7] Farneback, G. (2003). Two-frame motion estimation based on polynomial expansion. In *Scandinavian Conference on Image Analysis* (pp. 363–370). Springer, Berlin, Heidelberg.
- [8] Rother, C., Kolmogorov, V., & Blake, A. (2004). GrabCut: Interactive foreground extraction using iterated graph cuts. *ACM Transactions on Graphics*, 23(3), 309–314.
- [9] Vincent, L., & Soille, P. (1991). Watersheds in digital spaces: An efficient algorithm based on immersion simulations. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 13(6), 583–598.
- [10] Peng, J., Lu, J., & Zhang, D. (2020). Deep learning for precision agriculture: A survey. *Computers and Electronics in Agriculture*, 103, 105055.

Appendix A: Pipeline Module Summary

For rapid reference, Table A1 below summarises all ten pipeline modules with their implementation function names, output dimensionality, and computational complexity class.

Module	Function	Output Dim	Complexity
T1 Spectral	t1_multispectral(roi)	6 features	$O(N)$ pixels
T2 Texture	t2_glcmm_texture(roi)	5 features	$O(N)$ vectorised
T3 Depth	t3_depth_defocus(roi)	3 features	$O(N)$ Laplacian
T4 Flow	t4_optical_flow(prev, curr)	3 features	$O(N \cdot \text{iter})$ Farneback
T5 HSV	t5_hsv_segment(hsv)	4 binary masks	$O(N)$ threshold
T6 Morphology	t6_morphology(mask)	1 binary mask	$O(N \cdot k^2)$ conv
T7 Watershed	t7_watershed(bgr, mask)	Label array	$O(N \log N)$
T8 CNN	cnn_classify(patch)	$P \in [0,1]$	$O(\text{params})$ GPU
T9 RF	HARVESTER.predict(roi)	$P \in [0,1]$	$O(\text{trees} \cdot \text{depth})$

Module	Function	Output Dim	Complexity
T10 Arduino	serial.write(json)	Hardware cmd	O(1) serial

Table A1. Summary of all ten pipeline modules, their callable functions, output shapes, and computational complexity.

Appendix B: Ensemble Weight Sensitivity

The ensemble formula $\text{score} = w_{\text{RF}} \cdot \text{RF} + w_{\text{CNN}} \cdot \text{CNN} + w_{\text{HSV}} \cdot \text{HSV} + w_{\text{Depth}} \cdot \text{Depth}$ is parameterised by four weights constrained to sum to 1.0. The default configuration $w = (0.35, 0.35, 0.20, 0.10)$ was selected based on empirical cross-validation over a grid of weight combinations on the synthetic validation set. Sensitivity analysis (not shown) indicates that detection performance is relatively robust to perturbations of ± 0.05 in any individual weight, provided the RF and CNN components remain co-dominant (each ≥ 0.25). Reducing either below 0.20 degrades precision due to the higher false-positive rate of HSV and depth cues individually.